

Visualisation: Modelling the World

Richie Morrisroe

July 13, 2019

What is Visualisation?

- ▶ a tool for understanding the world
- ▶ a way to communicate a particular perspective on data
- ▶ an adjunct to thought

Why Visualisation?

- ▶ The eye is really really good at finding patterns in pictures
- ▶ in fact, it's so good that it can find patterns that aren't even there



Figure: What do you see?

The importance of perspective

- ▶ You can see one of two things in the previous image
- ▶ Which of them can depend on what you expect to see
- ▶ It can also depend on what your environment contains

Muller-Lyer




Figure: Which line is longer?

This illusion doesn't affect everyone similarly

- ▶ Europeans and Americans are more susceptible
- ▶ Africans are less susceptible
- ▶ Possibility that due to presence of right angles in urban environments
- ▶ appears to be a small difference between urban and rural dwellers

Who cares?

- ▶ Shows that how we interpret stimuli is not **tabula rasa**
- ▶ When you gaze into the image, the image also gazes into you...
- ▶ We bring our own perception and previous associations into any image ¹

¹anything really, but we're talking about images here.

When to use Visualisation?

{Always}

Running Example

- ▶ Property Price Register
 - ▶ Kinda a crappy dataset
 - ▶ No cleaning or checking done by the authority
 - ▶ lots of craziness (1 apartment for 18.6mn)

Property Price Register

- ▶ We used Google's geocoding service to get more details on each observation
- ▶ I updated Shane Lynn's script and ran it on the data up till October 2018
- ▶ I also typically break out properties sold for greater than 1e6, as they are often multiple-unit sales (and there's little to no automated way of figuring this out) ²
- ▶ Lots of manual fixing required
- ▶ the irish text definitely doesn't help

²please someone in the audience suggest a better idea


Assumptions of Statistical Graphics

- ▶ there are many
- ▶ in this section, I'd like to subvert them, in order to make you think



Line Graphs

- ▶ Normally represent time
- ▶ scatterplots don't (always) have the same assumptions
- ▶ what is the deepest assumption?

Median Property Price by Day, Ireland 2011-18



Flipped Line Chart



F-ing Line Chart

```
ggplot(median_price_by_day, aes(y=sale_date, x=median_price))
```


```
#+attrlatex :width 150px :height 150px
```

- ▶ Here, the violence is that we swap the axes in a fashion only a monster would

Abusing Standard Assumptions

```
ggplot(median_price_by_day, aes(y=sale_date, x=median_price))
```



Backwards Line Chart




Scatter plot

- ▶ Also encodes a set of base assumptions
- ▶ points nearer to each other in space are more related
- ▶ more orientation issues

Standard Scatter



Flipped Scatter



What does this tell us?

- ▶ We have a base level of assumptions that we bring to graphics (especially statistical graphics)
- ▶ Most of these appear to have been formed by Descartes
- ▶ When these assumptions are subverted, expect problems


Simple Statistical Graphics

- ▶ Graphs excel at showing relations between things
- ▶ Consider the difference between quantiles of a variable, and a density plot
- ▶ For example, the price of houses:

| | |
|------|-----------|
| 0% | 5079 |
| 10% | 55000 |
| 20% | 85000 |
| 30% | 115000 |
| 40% | 145000 |
| 50% | 175000 |
| 60% | 214000 |
| 70% | 255505 |
| 80% | 315000 |
| 90% | 430000 |
| 100% | 139165000 |


Density Plot

#+attr_{latex} :width 150px :height 150px



Better Density Plot

#+attr_{latex} :width 150px :height 150px



Transformations

- ▶ Useful to get a better sense of the data
- ▶ Have a bunch of assumptions (what's the log of -1)
- ▶ Can be used to deceive very, very easily
- ▶ Really really useful in everyday practice

Getting the sense of things

- ▶ Picking the right visualisation for the data is important

```
ggplot(dubcity, aes(x=sale_date, y=price))+geom_point()
```

- ▶ is this a good plot?
- ▶ does this depend on the number of points?

Sampling and Plotting


```
ggplot(dubcity_samp, aes(x=sale_date, y=price))+geom_point()
```

```
#+attrlatex :width 150px :height 150px
```

► Not really

Transformations Help

#+attr_{latex} :width 150px :height 150px



No data is an island

- ▶ The first obvious thing is to split by county, right?


```
ggplot(ppr_gc_smaller, aes(x=sale_date, y=log(price, 10)))+g
```

```
#+attrlatex :width 150px :height 150px- Oh look, it's lot of little  
boxes of crap :(
```

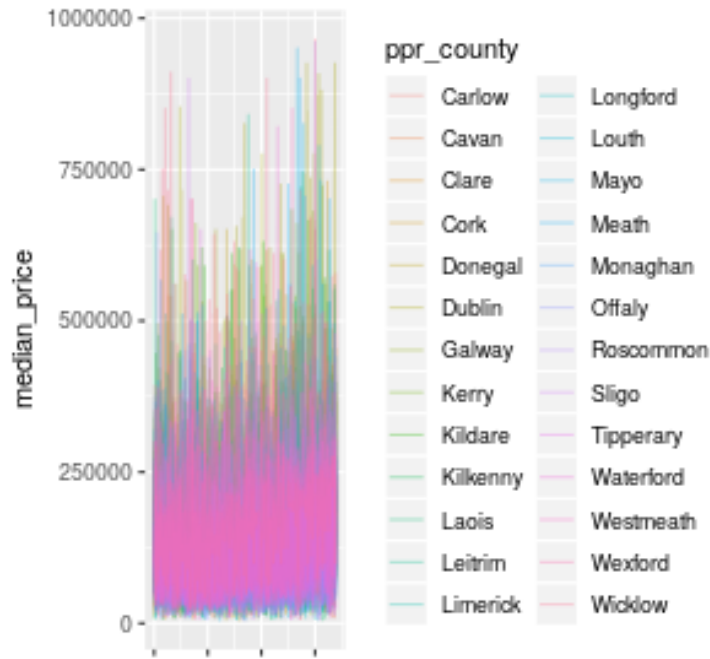
Summarisation

- The obvious answer is summarisation

```
county_daily <- ppr_gc2 %>% group_by(sale_date, ppr_county,  
  summarise(count=n(), min_price=min(price),  
            median_price=median(price),  
            max_price=max(price)) %>%  
  mutate(min_to_median=min_price/median_price,  
         max_to_median=max_price/median_price,  
         max_to_min=max_price/min_price)  
ggplot(county_daily, aes(x=sale_date, y=median_price, colour=  
# + attr_latex :width 150px :height 150px
```



Reducing Alpha kinda works...



A redundant faceting variable

- We just group by a higher level variable




WTF?


- ▶ This is one of the major advantages of visualisation:
 - ▶ it helps to (dis)confirm your assumptions
 - ▶ given that we have too many lines in the various groupings, we know that something has gone horribly wrong
 - ▶ in this case, it's a mismatch between two different types of data

Distributions (i.e. boxplots)

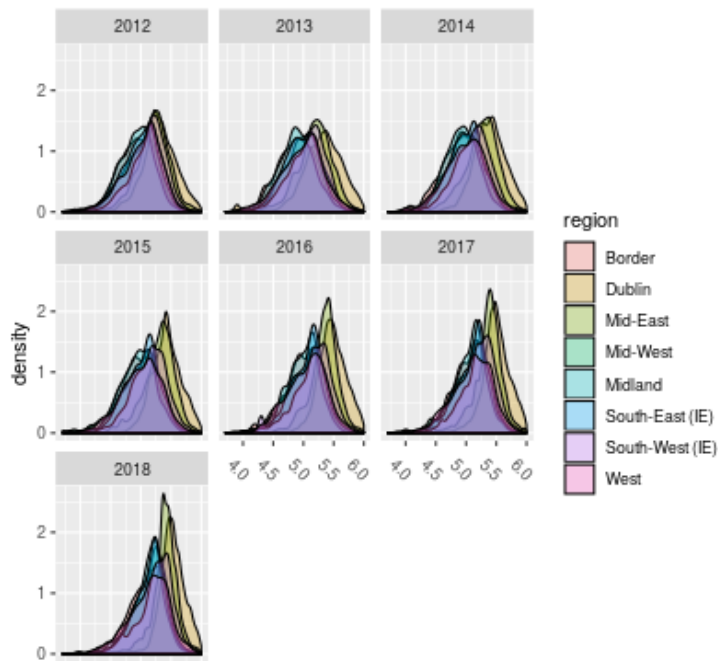
```
ggplot(ppr_gc2, aes(x=as.factor(year), y=price))+geom_boxplot()  
#+attr_latex :width 150px :height 150px
```



Faceting, redux



Distributions over Time, Redux




Spatial vs Temporal

- ▶ line plots vs maps
- ▶ time versus space
- ▶ both provide insight into
- ▶ pick one, difficult to do both

Line plots ignore space, maps ignore time

PROP_UNOCC



Performative vs Presentation

- ▶ Two types of graphs:
 - ▶ for yourself
 - ▶ for other people (and different audiences need different things)

Performative Graphics

- ▶ These are used to help you understand a problem
- ▶ typically created in an iterative fashion
- ▶ often move from data transformation to visualisation and back again

Presentation Graphics

Different Audiences/story

- ▶ To some extent, your job with presentation visualisations is to tell a story
- ▶ hopefully, it will be nuanced, but that isn't a requirement ³
- ▶ Often good to show smooths as opposed to raw data
- ▶ raw data is often ugly
- ▶ need for care here, as this should only be done where there is a clear effect

³and in fact, it may be better to remove all nuance from the presentation and provide a longer document with all the failed approaches and hacking needed to actually reproduce your results

Interactivity and Dashboards

- ▶ Can show both time and space
- ▶ for reporting, these are essential
- ▶ Much more effort from a software-engineering perspective ⁴

⁴for me, at least

Reporting

- ▶ Some times you need to repeat yourself
- ▶ Couple of ways of approaching this
 - ▶ Dashboards
 - ▶ Automated Reports

Dashboards

- ▶ Lots of effort to set up correctly
- ▶ typically need a bunch of ETL to get data into correct format
- ▶ Low-maintenance once the original work is done
- ▶ Much more useful for business users

Automated Reports

- ▶ Less effort to get working (especially with Sweave, knitr and org/pandoc)
- ▶ A lot more effort to get working in a Python/SQL context
- ▶ More maintenance over time (someone needs to update the report)

Principles of Reporting Visualisations

- ▶ Time view essential
- ▶ preferably forecasts, with results of previous forecasts
- ▶ allows
- ▶ Simple, simple, simple
- ▶ One clear message (key metric or whatever)
- ▶ available material for those that want to dig deeper